

SWE-534 NATURAL LANGUAGE PROCESSING PRACTICAL FILE

*Submitted towards the partial fulfilment
of the requirements of the award of the degree of*

Master of Technology In Software Engineering

Submitted To:-

Dr. Sanjay Patidar
Assistant Professor
Department of Software Engineering

Submitted By:-

Aman Chauhan (24/SWE/08)
II SEM, I YEAR



Delhi Technological University

(FORMERLY Delhi College of Engineering)

Bawana Road, New Delhi-110042

April, 2025

INDEX

S.No.	EXPERIMENT NAME	PAGE NO.	DATE	SIGNATURE
1	Build a complete text preprocessing pipeline to clean and prepare raw text for NLP models.	1	08/01/25	
2	Classify customer feedback as positive or negative using traditional machine learning models.	4	15/01/25	
3	Extract entities like company names, dates, and invoice amounts from text	8	22/01/25	
4	Classify news articles into predefined categories such as politics, sports, and technology.	12	29/01/25	
5	Generate concise summaries of lengthy legal documents using extractive and abstractive methods.	15	19/02/25	
6	Translate text from one language to another using neural machine translation models.	19	12/03/25	
7	Perform sentiment analysis on tweets or social media posts using a pre trained BERT model.	22	19/03/25	
8	Build a system that can answer questions based on a given context.	27	26/03/25	
9	Develop a chatbot that can handle basic customer queries.	31	02/04/25	
10	Generate human-like text for content creation using GPT models.	35	09/04/25	

EXPERIMENT - 1

AIM:

Build a complete text preprocessing pipeline to clean and prepare raw text for NLP models.

DESCRIPTION:

This experiment involves constructing a modular pipeline to transform raw, unstructured text (e.g., product reviews) into clean, standardized tokens suitable for NLP models. The pipeline sequentially applies:

1. **Tokenization:** Splitting text into words or subword units using SpaCy's efficient tokenizer.
2. **Lowercasing:** Normalizing text to lowercase to reduce vocabulary redundancy.
3. **Stopword Removal:** Filtering out common non-informative words (e.g., "the," "and") using NLTK's predefined stopwords lists.
4. **Lemmatization:** Reducing words to their base forms (e.g., "running" → "run") via SpaCy's lemmatizer, which considers part-of-speech tags for accuracy.

The output is a cleaned corpus ready for vectorization (e.g., TF-IDF, word embeddings) and downstream tasks like sentiment classification.

INTRODUCTION:

Effective text preprocessing is critical in NLP, as raw text often contains noise (typos, slang, punctuation) that hinders model performance. For businesses analyzing customer feedback, preprocessing ensures that sentiment analysis models focus on semantically meaningful content. For example, lemmatizing "disappointed" and "disappoints" to a common root ("disappoint") helps models generalize patterns. By standardizing text, this pipeline enhances feature extraction and improves the reliability of insights derived from user-generated content.

IMPLEMENTATION:

```
from google.colab import files
import pandas as pd
import nltk
import spacy
from nltk.corpus import stopwords
```

```
# Download NLTK stopwords
nltk.download('stopwords')
```

```
# Load SpaCy model for lemmatization
nlp = spacy.load("en_core_web_sm")
```

```
# Step 1: Load sample text data
```

```
url = "https://raw.githubusercontent.com/AmanChn/Dataset/refs/heads/main/amazon_reviews.csv"
```

```

data = pd.read_csv(url)

# Display the first few rows of the dataset
print("Original Data:")
print(data.head())

# Step 2: Define preprocessing functions
def preprocess_text(text):
    # Handle missing or non-string values
    if not isinstance(text, str):
        return ""

    # Convert to lowercase
    text = text.lower()

    # Tokenization and stopword removal
    tokens = [word for word in text.split() if word not in stopwords.words('english')]

    # Lemmatization using SpaCy
    doc = nlp(" ".join(tokens))
    lemmatized_tokens = [token.lemma_ for token in doc]

    # Return the cleaned text
    return " ".join(lemmatized_tokens)

# Apply preprocessing to the reviews column
# Replace 'review' with the correct column name if it's different
data['cleaned_review'] = data['reviewText'].apply(preprocess_text)

# Display the cleaned text
print("\nCleaned Data:")
print(data[['reviewText', 'cleaned_review']].head())

# Save cleaned data for further analysis
data.to_csv("cleaned_amazon_reviews.csv", index=False)
files.download("cleaned_amazon_reviews.csv")

```

OUTPUT:

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
⇒ Original Data:
  Unnamed: 0  reviewerName  overall  \
0           0           NaN      4.0
1           1           0mie      5.0
2           2           1K3      4.0
3           3           1m2      5.0
4           4  2&1/2Men      5.0

          reviewText  reviewTime  day_diff  \
0                No issues.  2014-07-23    138
1  Purchased this for my device, it worked as adv...  2013-10-25    409
2  it works as expected. I should have sprung for...  2012-12-23    715
3  This think has worked out great.Had a diff. br...  2013-11-21    382
4  Bought it with Retail Packaging, arrived legit...  2013-07-13    513

  helpful_yes  helpful_no  total_vote  score_pos_neg_diff  \
0           0           0           0           0
1           0           0           0           0
2           0           0           0           0
3           0           0           0           0
4           0           0           0           0

  score_average_rating  wilson_lower_bound
0                   0.0                   0.0
1                   0.0                   0.0
2                   0.0                   0.0
3                   0.0                   0.0
4                   0.0                   0.0

```

Cleaned Data:

```

          reviewText  \
0                No issues.
1  Purchased this for my device, it worked as adv...
2  it works as expected. I should have sprung for...
3  This think has worked out great.Had a diff. br...
4  Bought it with Retail Packaging, arrived legit...

  cleaned_review
0          issue .
1  purchase device , worked advertise . never muc...
2  work expect . spring high capacity . think mak...
3  think work great.had diff . bran 64 gb card go...
4  buy retail packaging , arrive legit , orange e...

```

EXPERIMENT - 2

AIM:

Classify customer feedback as positive or negative using traditional machine learning models.

DESCRIPTION:

Sentiment Classification of Customer Reviews focuses on training a machine learning model to automatically detect sentiment (positive/negative) in customer feedback. The workflow includes:

1. **Data Preparation:** Using labeled review datasets (e.g., Amazon product reviews) and preprocessing them via the text pipeline from Experiment 1.
2. **Feature Engineering:** Converting cleaned text into numerical vectors using TF-IDF, which highlights terms that are frequent in a document but rare across the corpus.
3. **Model Training:** Implementing a Logistic Regression classifier (Scikit-learn) due to its efficiency with high-dimensional text data and interpretable coefficients.
4. **Evaluation:** Assessing performance with accuracy (overall correctness) and F1-score (balance of precision and recall), which is critical for imbalanced datasets.

The output is a deployable model that identifies satisfaction trends, enabling businesses to prioritize product improvements.

INTRODUCTION:

Sentiment analysis automates the extraction of actionable insights from unstructured customer feedback, a task impractical to perform manually at scale. Traditional models like Logistic Regression remain popular for their transparency and speed, especially when businesses need to trace why a review was classified as negative (e.g., identifying frequent terms like “defective” or “slow”). By combining TF-IDF—which weights words by their diagnostic relevance—with a robust classifier, this pipeline transforms subjective opinions into quantifiable metrics. For instance, detecting spikes in negative sentiment linked to specific product features allows companies to allocate resources strategically, turning raw data into competitive advantage.

IMPLEMENTATION:

```
# Install dependencies as needed:
# pip install kagglehub[pandas-datasets]
from google.colab import files
import kagglehub
from kagglehub import KaggleDatasetAdapter
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, f1_score, classification_report
```

```

import nltk
import spacy
from nltk.corpus import stopwords

# Download NLTK stopwords
nltk.download('stopwords')

# Load SpaCy model for lemmatization
nlp = spacy.load("en_core_web_sm")

# Step 1: Load labeled customer review data from Kaggle
file_path = "Dataset-SA.csv" # Set to the specific file in the dataset
df = kagglehub.load_dataset(
    KaggleDatasetAdapter.PANDAS,
    "niraliivaghani/flipkart-product-customer-reviews-dataset",
    file_path,
)

print("First 5 records from the dataset:")
print(df.head())

# Check and preprocess the dataset
# Ensure dataset has 'Review' and 'Sentiment' columns
df = df.dropna(subset=['Review', 'Sentiment']) # Remove rows with missing values

# Step 2: Preprocess the data using the pipeline from Experiment 1
def preprocess_text(text):
    if not isinstance(text, str):
        return ""
    text = text.lower()
    tokens = [word for word in text.split() if word not in stopwords.words('english')]
    doc = nlp(" ".join(tokens))
    lemmatized_tokens = [token.lemma_ for token in doc]
    return " ".join(lemmatized_tokens)

df['cleaned_review'] = df['Review'].apply(preprocess_text)

# Step 3: Extract features using TF-IDF
tfidf = TfidfVectorizer(max_features=5000)
X = tfidf.fit_transform(df['cleaned_review']).toarray()
y = df['Sentiment'].apply(lambda x: 1 if x.lower() == 'positive' else 0) # Encode sentiment

# Step 4: Train a Logistic Regression model for classification
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

model = LogisticRegression()
model.fit(X_train, y_train)

# Step 5: Evaluate using metrics like accuracy and F1-score
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("\nEvaluation Metrics:")
print(f'Accuracy: {accuracy:.2f}')
print(f'F1 Score: {f1:.2f}')
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Save the processed data and model
df.to_csv("classified_reviews_kaggle.csv", index=False)
files.download("classified_reviews_kaggle.csv")

```

OUTPUT:

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
<ipython-input-1-593ea656cf36>:25: DeprecationWarning: load_dataset is deprecated and will be removed in future version.
df = kagglehub.load_dataset(
Downloading from https://www.kaggle.com/api/v1/datasets/download/niraliivaghani/flipkart-product-customer-reviews-dataset?dataset_version_number=1&file_name=Data
100%|██████████| 3.79M/3.79M [00:00<00:00, 56.8MB/s]Extracting zip of Dataset-SA.csv...

First 5 records from the dataset:
   product_name  product_price  Rate \
0  Candes 12 L Room/Personal Air Cooler?????(Whi...      3999      5
1  Candes 12 L Room/Personal Air Cooler?????(Whi...      3999      5
2  Candes 12 L Room/Personal Air Cooler?????(Whi...      3999      3
3  Candes 12 L Room/Personal Air Cooler?????(Whi...      3999      1
4  Candes 12 L Room/Personal Air Cooler?????(Whi...      3999      3

   Review                                     Summary \
0  super! great cooler excellent air flow and for this p...
1  awesome best budget 2 fit cooler nice cooling
2  fair the quality is good but the power of air is de...
3  useless product very bad product its a only a fan
4  fair ok ok product

Sentiment
0 positive
1 positive
2 positive
3 negative
4 neutral

```


Evaluation Metrics:

Accuracy: 0.92

F1 Score: 0.95

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.60	0.73	6496
1	0.92	0.99	0.95	29582
accuracy			0.92	36078
macro avg	0.92	0.79	0.84	36078
weighted avg	0.92	0.92	0.91	36078

EXPERIMENT - 3

AIM:

Extract entities like company names, dates, and invoice amounts from text.

DESCRIPTION:

This experiment involves extracting structured information (e.g., vendor names, dates, amounts) from unstructured invoice documents using NLP. The pipeline includes:

1. **Model Selection:** Leveraging SpaCy's pretrained NER model (e.g., en_core_web_sm) or Hugging Face's transformer-based models (e.g., BERT) for high-accuracy entity detection.
2. **Fine-Tuning (Optional):** Adapting the model to domain-specific terminology (e.g., invoice IDs, tax codes) using labeled invoice datasets if the pretrained model underperforms.
3. **Data Structuring:** Exporting detected entities (e.g., DATE, ORG, MONEY) into a CSV file via Pandas, enabling integration with accounting systems.

The output automates data entry, reducing manual errors and processing time for financial workflows.

INTRODUCTION:

Manual invoice processing is prone to human error and inefficiency, especially for businesses handling hundreds of invoices monthly. NER addresses this by transforming unstructured text into machine-readable fields. Pretrained models like SpaCy's pipeline identify common entities out-of-the-box, but fine-tuning may be necessary to capture industry-specific terms (e.g., "PO number" or "GSTIN"). For example, accurately extracting "\$1,500" as an invoice amount and associating it with the correct "Due Date" ensures timely payments. By automating this process, companies can reallocate resources to strategic tasks while maintaining audit-ready records.

IMPLEMENTATION:

```
# Install required libraries
!pip install spacy pandas
!python -m spacy download en_core_web_sm # Download SpaCy's small English model
```

```
# Import libraries
import spacy
import pandas as pd
from google.colab import files
import json
import re
```

```
# Upload your Excel file to Colab
uploaded = files.upload() # This will prompt you to upload your Excel file
excel_file = list(uploaded.keys())[0] # Get the uploaded file name
```

```

# Load the Excel dataset
df = pd.read_excel(excel_file)
print("Dataset loaded successfully:")
print(df.head())

# Load SpaCy's pretrained NER model
nlp = spacy.load("en_core_web_sm")

# Function to extract entities from invoice text
def extract_entities(text):
    doc = nlp(text)
    entities = {}

    # Define patterns for specific invoice-related entities
    total_amount_pattern = r"\$[\d,]+\.\d*" # Matches amounts like $1000 or $550.50
    date_pattern = r"\d{1,2} \w+ \d{4}" # Matches dates like "5 July 2025"

    # Use SpaCy NER for general entities (e.g., organizations, dates)
    for ent in doc.ents:
        if ent.label_ == "ORG":
            entities["COMPANY_NAME"] = ent.text
        elif ent.label_ == "DATE":
            entities["PAYMENT_DATE"] = ent.text

    # Use regex for specific invoice fields
    total_amount = re.search(total_amount_pattern, text)
    if total_amount:
        entities["TOTAL_AMOUNT"] = total_amount.group()

    # Parse the Final_Output JSON to merge with extracted entities
    try:
        final_output = json.loads(df.loc[df['Input'] == text, 'Final_Output'].values[0])
        entities.update(final_output) # Merge with predefined output
    except:
        pass

    return entities

# Apply entity extraction to the dataset
df['Extracted_Entities'] = df['Input'].apply(extract_entities)

# Convert extracted entities to a structured DataFrame
entity_df = pd.json_normalize(df['Extracted_Entities'])

```

```

# Combine original input with extracted entities
result_df = pd.concat([df[['Input']], entity_df], axis=1)

# Display the results
print("\nExtracted Entities:")
print(result_df.head())

# Save the results to a CSV file
result_df.to_csv('invoice_ner_results.csv', index=False)
print("\nResults saved to 'invoice_ner_results.csv'")

# Download the CSV file
files.download('invoice_ner_results.csv')

# Optional: Fine-tuning SpaCy (uncomment and modify if you have labeled data)
"""
# Example for fine-tuning (requires labeled training data)
from spacy.training.example import Example

# Sample training data (you'd replace this with your labeled dataset)
TRAIN_DATA = [
    ("Cream and White Simple Minimalist Catering Services Invoice", {
        "entities": [(0, 54, "COMPANY_NAME"), (55, 60, "TOTAL_AMOUNT")]
    })
]

# Fine-tuning code
def train_spacy_model(nlp, train_data, iterations=10):
    optimizer = nlp.resume_training()
    for i in range(iterations):
        for text, annotations in train_data:
            doc = nlp.make_doc(text)
            example = Example.from_dict(doc, annotations)
            nlp.update([example], drop=0.5, sgd=optimizer)
    return nlp

# Train the model
nlp = train_spacy_model(nlp, TRAIN_DATA)
nlp.to_disk("fine_tuned_ner_model") # Save the fine-tuned model
"""

```

OUTPUT:

```
✓ Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')
⚠ Restart to reload dependencies
If you are in a Jupyter or Colab notebook, you may need to restart Python in
order to load all the package's dependencies. You can do this by selecting the
'Restart kernel' or 'Restart runtime' option.
Choose Files converted i...dataset.xlsx
• converted_invoice_dataset.xlsx(application/vnd.openxmlformats-officedocument.spreadsheetml.sheet) - 24062 bytes, last modified: 8/22/2024 - 100% done
Saving converted_invoice_dataset.xlsx to converted_invoice_dataset.xlsx
Dataset loaded successfully:

Input \
0 Cream and White Simple Minimalist Catering Ser...
1 Beige Elegant Professional Business Invoice\n\...
2 Black and White Clean Modern Invoice\n\nConsul...
3 Black and White Minimalist Business Invoice\n\...
4 White Minimalist Business Invoice\n\nSUBTOTALN...

Final Output
0 {"TOTAL_AMOUNT": "$1000", "DUE_AMOUNT": "$550"...
1 {"INVOICE_NUMBER": "#01234", "BILLED_TO": "Est...
2 {"BILL_TO": "SALFORD & CO.", "BANK_NAME": "Bor...
3 {"INVOICE_NUMBER": "12345", "BILLED_TO": "Marc...
4 {"INVOICE_NUMBER": "#123456", "DATE_ISSUED": "...
```

```
Extracted Entities:

Input \
0 Cream and White Simple Minimalist Catering Ser...
1 Beige Elegant Professional Business Invoice\n\...
2 Black and White Clean Modern Invoice\n\nConsul...
3 Black and White Minimalist Business Invoice\n\...
4 White Minimalist Business Invoice\n\nSUBTOTALN...

COMPANY_NAME PAYMENT_DATE \
0 Borcelle Bank\nAccount Name 5 July 2025
1 Borcelle Bank\nAccount 0123 4567 8901
2 Borcelle Bank\nName 30 days
3 White Minimalist Business Invoice 4567
4 White Minimalist Business Invoice\n\nSUBTOTALNO 12345

TOTAL_AMOUNT DUE_AMOUNT INVOICE_NUMBER ITEM_DESCRIPTION QTY \
0 $1000 $550 #612345 Grilled Chicken 2
1 $195.00 NaN #01234 60-minute full body massage NaN
2 $6,875 NaN INV-01234 NaN NaN
3 $20.00 NaN 12345 Architecture NaN
4 NaN NaN #123456 Copywriting 5

UNIT_PRICE BANK_NAME ... FREELANCER COMPANY PAYPAL PAY_BY \
0 $200 Borcelle Bank ... NaN NaN NaN NaN
1 NaN Borcelle Bank ... NaN NaN NaN NaN
2 NaN Borcelle Bank ... NaN NaN NaN NaN
3 NaN Really Great Bank ... NaN NaN NaN NaN
4 NaN Fauget ... NaN NaN NaN NaN
```

```
BANK_CODE ADMINISTRATOR BANK_ACCOUNT DUE_DATE SHIPPING SERVICE
0 NaN NaN NaN NaN NaN NaN
1 NaN NaN NaN NaN NaN NaN
2 NaN NaN NaN NaN NaN NaN
3 NaN NaN NaN NaN NaN NaN
4 NaN NaN NaN NaN NaN NaN

[5 rows x 64 columns]

Results saved to 'invoice_ner_results.csv'
'\n# Example for fine-tuning (requires labeled training data)\nfrom spacy.training.example import Example\n\n# Sample training data (you'd replace this with your labeled dataset)\nTRAIN_DATA = [\n    ("Cream and White Simple Minimalist Catering Services Invoice", {\n        "entities": [(0, 54, "COMPANY_NAME"), (55, 60, "TOTAL_AMOUNT")]\n    })\n]\n\n# Fine-tuning code\ndef train_spacy_model(nlp, train_data, iterations=10):\n    optimizer = nlp.resume_training()\n    for i in range(iterations):\n        for text, annotations in train_data:\n            doc = nlp.make_doc(text)\n            example = Example.from_dict(doc, annotations)\n            nlp.update([example], drop=0.5, sgd=optimizer)\n    return nlp\n\n# Train the model\nnlp = train_spacy_model(nlp, TRAIN_DATA)\nnlp.to_disk("fine_tuned_ner_model") # Save the fine-tuned model\n
```

EXPERIMENT - 4

AIM:

Classify news articles into predefined categories such as politics, sports, and technology.

DESCRIPTION:

This experiment automates the categorization of news articles into predefined topics (e.g., politics, sports, technology) using machine learning. The workflow includes:

1. Data Preparation: Aggregating labeled news datasets (e.g., BBC News) and preprocessing text using tokenization, stopwords removal, and lemmatization
2. Feature Extraction:
 - TF-IDF: Highlighting discriminative terms (e.g., "election" for politics, "goal" for sports).
 - Word2Vec (via Gensim): Capturing semantic relationships through word embeddings.
3. Model Training:
 - Random Forest: Handling non-linear decision boundaries and high-dimensional features.
 - Naive Bayes: Offering fast training for real-time applications.
4. Evaluation: Reporting precision (minimizing misclassified articles), recall (capturing all relevant articles), and F1-score (balancing both).

The output is a classifier that organizes articles into topics, enabling personalized user recommendations.

INTRODUCTION:

News platforms generate vast volumes of content daily, making manual categorization impractical. Automated topic classification streamlines content discovery, ensuring users receive relevant articles (e.g., tech enthusiasts see AI advancements, not sports scores). TF-IDF prioritizes domain-specific keywords, while Word2Vec captures context (e.g., "court" in sports vs. legal news). For businesses, accurate classification enhances user engagement by tailoring feeds to individual interests. For instance, a misclassified "SpaceX launch" under "entertainment" could frustrate readers, whereas precise tagging improves retention. By deploying this model, media companies can scale content delivery while maintaining editorial coherence.

IMPLEMENTATION:

```
!pip install scipy==1.11.4 scikit-learn==1.3.2 gensim==4.3.2 pandas==2.0.3 kaggle==1.6.12
```

```
# Import libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
```

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report
from google.colab import files
import os

# Option 1: Upload dataset manually
print("Please upload your AG News dataset (e.g., training_data.csv from Kaggle)")
uploaded = files.upload()
dataset_file = list(uploaded.keys())[0]
train_dataset = pd.read_csv(dataset_file)

# Display dataset info
print("Dataset loaded successfully:")
print(train_dataset.head())
print("\nDataset shape:", train_dataset.shape)

# Assuming columns are 'text' and 'label'
X = train_dataset['text']
y = train_dataset['label']

# Preprocess and split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Feature extraction using TF-IDF
tfidf_vectorizer = TfidfVectorizer(max_features=5000, stop_words='english')
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
X_test_tfidf = tfidf_vectorizer.transform(X_test)

# Train Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
rf_classifier.fit(X_train_tfidf, y_train)

# Predict and evaluate Random Forest
rf_predictions = rf_classifier.predict(X_test_tfidf)
print("\nRandom Forest Evaluation:")
print("Precision:", precision_score(y_test, rf_predictions, average='weighted'))
print("Recall:", recall_score(y_test, rf_predictions, average='weighted'))
print("F1-Score:", f1_score(y_test, rf_predictions, average='weighted'))
print("\nDetailed Classification Report:")
print(classification_report(y_test, rf_predictions))

# Train Naive Bayes Classifier
nb_classifier = MultinomialNB()

```

```

nb_classifier.fit(X_train_tfidf, y_train)

# Predict and evaluate Naive Bayes
nb_predictions = nb_classifier.predict(X_test_tfidf)
print("\nNaive Bayes Evaluation:")
print("Precision:", precision_score(y_test, nb_predictions, average='weighted'))
print("Recall:", recall_score(y_test, nb_predictions, average='weighted'))
print("F1-Score:", f1_score(y_test, nb_predictions, average='weighted'))
print("\nDetailed Classification Report:")
print(classification_report(y_test, nb_predictions))

# Save the results to a CSV file
results = pd.DataFrame({
    'Text': X_test,
    'Actual_Label': y_test,
    'RF_Predicted_Label': rf_predictions,
    'NB_Predicted_Label': nb_predictions
})
results.to_csv('news_topic_classification_results.csv', index=False)
print("\nResults saved to 'news_topic_classification_results.csv'")

# Download the results
files.download('news_topic_classification_results.csv')

```

OUTPUT:

```

Accuracy and cm of training set:
Confusion Matrix:
[[163193    0]
 [    0 163927]]
Accuracy Score 1.0
Accuracy and cm of test set:
Confusion Matrix:
[[40977    10]
 [   12    81]]
Accuracy Score 0.9994644595910419
Precision Score 0.8901098901098901
Recall Score 0.8709677419354839
F1 Score 0.8804347826086956
ROC AUC Score 0.9353618810685056

```


EXPERIMENT - 5

AIM:

Generate concise summaries of lengthy legal documents using extractive and abstractive methods.

DESCRIPTION:

This experiment compares two approaches to condense lengthy legal documents:

1. **Extractive Summarization (TextRank):** Uses Gensim's TextRank algorithm to identify and rank key sentences based on their similarity to others, preserving original wording.
2. **Abstractive Summarization (BART/T5):** Leverages Hugging Face's pretrained models (e.g., BART, T5) to generate paraphrased summaries that capture the document's essence in novel sentences.

The pipeline processes input text (e.g., contracts, court rulings), applies both methods, and evaluates summary quality using metrics like ROUGE (measuring overlap with human-written summaries) and human feedback for coherence.

INTRODUCTION:

Legal documents often span hundreds of pages filled with redundant clauses and complex jargon, making manual summarization labor-intensive and error-prone. Extractive methods like TextRank efficiently highlight critical passages (e.g., penalty clauses, obligations) but may lack contextual fluency. Abstractive models, while computationally heavier, mimic human-like summarization by rephrasing content (e.g., converting "The party shall remit payment within 30 days" to "Payment deadline: 30 days"). For law firms, automating this process accelerates case preparation—enabling lawyers to focus on strategy rather than skimming documents. Comparing both methods allows organizations to choose speed (extractive) versus nuance (abstractive) based on their needs.

IMPLEMENTATION:

```
!pip install datasets transformers evaluate
```

```
!pip install rouge_score
```

```
# Import libraries
```

```
from datasets import load_dataset
```

```
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM
```

```
import evaluate
```

```
# Load the Multi-LexSum dataset
```

```
multi_lexsum = load_dataset("allenai/multi_lexsum", name="v20220616")
```

```
# Access an example from the validation split
```

```
example = multi_lexsum["validation"][0]
```

```

source_documents = example["sources"] # List of source document texts for the case
long_summary = example["summary/long"] # Multi-paragraph summary
short_summary = example["summary/short"] # One-paragraph summary
tiny_summary = example["summary/tiny"] # One-sentence summary

# Display the first source document and summaries
print("Source Document:", source_documents[0])
print("\nLong Summary:", long_summary)
print("\nShort Summary:", short_summary)
print("\nTiny Summary:", tiny_summary)

# Load the tokenizer and model for summarization
tokenizer = AutoTokenizer.from_pretrained("facebook/bart-large-cnn")
model = AutoModelForSeq2SeqLM.from_pretrained("facebook/bart-large-cnn")

# Function for summarization
def summarize(text):
    inputs = tokenizer.encode("summarize: " + text, return_tensors="pt", max_length=1024,
truncation=True)
    summary_ids = model.generate(inputs, max_length=130, min_length=30,
length_penalty=2.0, num_beams=4, early_stopping=True)
    return tokenizer.decode(summary_ids[0], skip_special_tokens=True)

# Generate summaries
abstractive_summary = summarize(source_documents[0])

# Display summaries
print("\nAbstractive Summary:", abstractive_summary)

# Evaluate summaries using ROUGE metric
rouge = evaluate.load("rouge")
reference = long_summary # Use the long summary as the reference
candidate = abstractive_summary # Compare the abstractive summary
scores = rouge.compute(predictions=[candidate], references=[reference])

print("\nROUGE Scores:", scores)

```

OUTPUT:

Installing collected packages: rouge_score
Successfully installed rouge_score-0.1.2
Source Document: Page 1
LEXSEE 2003 U.S. DIST. CT. PLEADINGS 3030 View U.S. District Court Opinion View Original Source Image of This Document
SUSAN STOCKING, for herself and all other similarly situated, Plaintiff, v. AT&T CORP., Defendant.
Case No. 03-0421-CV-W-HFS UNITED STATES DISTRICT COURT FOR THE WESTERN DISTRICT OF MISSOURI, WESTERN DIVISION 2003 U.S. Dist. Ct. Pleadings 3030; 2004 U.S. Dist. Ct. Pleadings LEXIS 9181
January 23, 2004 Complaint
VIEW OTHER AVAILABLE CONTENT RELATED TO THIS DOCUMENT: U.S. District Court: Motion(s) COUNSEL: **[**1]** Sylvester "Sly" James, Jr. MO # 33617, Michael J. Mohlman
1. Plaintiff Susan Stocking is a citizen of the state of Missouri, was employed by Defendant, and has filed an EEOC complaint and obtained the right to sue. See
2. AT&T employs Plaintiff and thousands **[**2]** of other women like her. AT&T is a citizen of the states of New

2003 U.S. Dist. Ct. Pleadings 3030, *1; 2004 U.S. Dist. Ct. Pleadings LEXIS 9181, **2

Page 2

York where it is incorporated and of New Jersey where its principal place of business is located; and can be served with service of process at The Corporation
3. Aetna Insurance Company ("Health Insurer") was a health insurer for Plaintiff and/or AT&T.
[*2] 4. Venue is proper in this Court as Defendant does business in this judicial district and has offered or provided health insurance to its employees in this
5. This Court has subject matter jurisdiction based on federal question jurisdiction, 28 U.S.C. § 1331, and 42 U.S.C. § 2000e - 5(f)(3).
A. Summary of Claims
6. Prescription medication related to reproduction is routinely covered for men, but not for women. Prescription contraception, which is used only by women, is
7. AT&T's exclusion of prescription contraception has an adverse disparate impact on Ms. Stocking and other members of the proposed class. Because prescription
8. As a result of AT&T decision to exclude contraceptives from its benefits plan, Ms. Stocking and other members of the proposed class are being discriminated
B. Plaintiff Susan Stocking
9. Plaintiff Susan Stocking has been employed by AT&T on a full-time basis since June 12, 1995. As part of the terms and conditions of her employment, Ms. Stock
10. Ms. Stocking was, at all times relevant to this cause of action, a woman of childbearing age who was concerned with the prevention of an unwanted or unplanned
11. On August 26, 2002, Ms. Stocking filed a charge with the EEOC at its Kansas City Area Office **[**5]** in Kansas City, Kansas alleging that AT&T's failure to provide
12. As a result, Ms. Stocking received a decision and a right-to-sue letter from the EEOC (copies of which are attached hereto as "Exhibit A" and Exhibit B").

Page 3

[*4] 13. On November 29, 2002, the EEOC issued a "Determination" to Plaintiff Susan Stocking, in which it considered AT&T's benefit plan and concluded that:
"...it is the Commission's position that the pre-July 2002 exclusion violates both Title VII and the PDA, since the statutes cover prescription contraceptives and
14. AT&T's refusal to provide the same benefits to both men and women has caused an economic hardship on Ms. Stocking and other members of the proposed class that
II. AT&T's Health Plan
15. As terms and conditions of their employment, AT&T offered Ms. Stocking the opportunity to enroll in one of three health plans (Plan). The three options were
16. Regular full-time occupational employees are eligible for coverage at the company's expense on the first day of the month in which the employee attains six
[*5] 17. Ms. Stocking elected to take advantage of the PPO option.
18. The Plan also provides for a Prescription Drug Benefit Plan which was administered by Merck-Medco (Drug Plan). The Drug Plan provided drug benefits to participants
19. Despite covering other preventative medical services and prescriptions, neither the Plan nor the Drug Plan provided for prescription drugs and/or devices used
20. If plaintiff became pregnant, however, the Plan would have covered the costs of either an abortion or continuing the pregnancy to term - whichever she chose
III. Harm to Ms. Stocking and Other Class Members
21. As a direct and proximate result of the AT&T Plan's failure to cover contraception **[**8]** to prevent pregnancy, Ms. Stocking and other proposed members of the
22. On information and belief, AT&T employs hundreds of women of reproductive age who use prescription contraception.

2003 U.S. Dist. Ct. Pleadings 3030, *5; 2004 U.S. Dist. Ct. Pleadings LEXIS 9181, **8

Page 4

23. If contraception were treated on an equal basis with other prescriptions under the Plan and the Drug Plan, Ms. Stocking and other proposed class members would
[*6] IV. Factual Framework
24. For over thirty years of their lives, women have the biological potential for pregnancy. Contraception is a drug or device that prevents pregnancy. Without
25. The Food and Drug Administration (FDA) has approved five methods of reversible prescription contraception: oral contraception; Morplant; Depo-Provera; intra
26. Women bear all of the physical burdens of pregnancy, which are quite substantial. **[**10]** Pregnancy itself can put a woman's life at risk. Ectopic pregnancy
27. Pregnancy also poses non-life threatening health risks for women. The morbidity rate during pregnancy is quite high. Twenty-two percent of all pregnant women
28. The more pregnancies she bears, the greater the likelihood a woman will suffer one or more of the myriad life and/or health-threatening complications of pregnancy
29. For women with pre-existing medical conditions, even one pregnancy can pose grave health risks. Preexisting medical conditions that are exacerbated by pregnancy
30. Unintended pregnancy poses far greater health risks to women and children than does intended pregnancy. The medical risks of unintended pregnancy are well documented

2003 U.S. Dist. Ct. Pleadings 3030, *8; 2004 U.S. Dist. Ct. Pleadings LEXIS 9181, **12

Page 5

dangerous, and may even be deadly, for women with hypertension or diabetes. These conditions are best managed when medical care is begun before conception. In any
31. Unintended pregnancy is both frequent and widespread in the United States. Forty-nine percent of all pregnancies in the United States are unintended. Among
32. Contraception enables women to plan their pregnancies and time the spacing between pregnancies. The shorter the interval between her pregnancies, the greater
33. Furthermore, even in an otherwise healthy woman, pregnancy poses medical risks that are significantly greater than the risks of using contraception. In any
[*9] 34. Due to the wide variation in effectiveness, cost, and medical appropriateness of available forms of contraception, choice of contraceptive method is essential
35. Women with medical conditions **[**14]** that require pregnancy avoidance, in particular, require a full range of contraceptive options because their medical conditions
36. For all of the above reasons, a recent study by the Institute of Medicine recommends improving contraceptive coverage in health plans in order to reduce the
37. The physical burdens of pregnancy increase the risk of interruption to a woman's education, career and professional development opportunities. The ability
38. Inadequate insurance coverage of contraception has substantial adverse economic consequences for the 67% of American women of reproductive age who rely on

2003 U.S. Dist. Ct. Pleadings 3030, *10; 2004 U.S. Dist. Ct. Pleadings LEXIS 9181, **15

Page 6

V. Statutory Framework
39. Title VII provides that: "It shall be an unlawful employment practice for an employer to ... discriminate against any individual with respect to his compensation
40. In 1978, Congress enacted the Pregnancy Discrimination Act ("PDA") which provides **[**16]** that the term "because of sex" in Title VII includes, but is not limited to
41. Contraception is "pregnancy-related" within the meaning of the PDA because it is medical treatment that provides women with the ability to control their birth
42. The exclusion of contraception from the Health Plan also has an adverse disparate impact on women in violation of Title VII because it forces them either to
VI. Class Action Allegations
43. The proposed class of plaintiffs in this case consists of:
All female employees of AT&T Corp. covered or offered to be covered by health insurance who used prescription contraceptives from August 1995, to the present (hereinafter
44. This action is properly maintainable as a class action under Fed. R. Civ. P. 23(a). Plaintiff is informed and believes that the class is so numerous that joinder of all
45. Commonality is met. Plaintiff class members have common issues of law:
(i) whether the failure to provide coverage for female prescription contraceptives a violation of Title VII and/or the PDA by disparate treatment or disparate treatment
(ii) whether AT&T violated Title VII and/or PDA;
(iii) whether Health Insurer violated Title VII and/or PDA;
[*12] (iv) the measure of damages for Plaintiffs who obtain contraceptives anyway; and

EXPERIMENT - 6

AIM:

Translate text from one language to another using neural machine translation models.

DESCRIPTION:

This experiment implements a translation pipeline to convert text between languages (e.g., English to Spanish) using Hugging Face's MarianMT models. The workflow includes:

1. **Model Selection:** Deploying pretrained MarianMT models (e.g., Helsinki-NLP/opus-mt-en-es for English-Spanish) for immediate use.
2. **Fine-Tuning:** Adapting the model to domain-specific terminology (e.g., legal, medical) using custom bilingual datasets if higher accuracy is required.
3. **API Integration:** Wrapping the model in a REST API via FastAPI or Flask, enabling real-time translation within chatbots or customer service platforms.

The output is a scalable solution that breaks language barriers, allowing businesses to serve global audiences seamlessly.

INTRODUCTION:

As companies expand globally, language differences hinder customer interaction—65% of consumers prefer support in their native language. Pretrained models like MarianMT offer immediate multilingual capabilities but may struggle with niche vocabulary (e.g., technical jargon in banking). Fine-tuning on domain-specific data ensures terms like “APR” or “IBAN” are translated contextually. For instance, a Spanish user’s query about “tarjeta de crédito” (credit card) can be instantly converted to English for a monolingual agent, then responses translated back. By embedding this pipeline into APIs, businesses automate cross-lingual communication, enhancing satisfaction while reducing reliance on human translators.

IMPLEMENTATION:

```
# Save this as `translation_api.py`

# Import required libraries
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
from transformers import MarianTokenizer, MarianMTModel
import torch

# Define the FastAPI app
app = FastAPI(
    title="Machine Translation API",
    description="API for translating text between languages using MarianMT",
    version="1.0.0"
)
```

```

# Define the language pair and load the pretrained MarianMT model
src_lang = "en" # Source language: English
tgt_lang = "es" # Target language: Spanish
model_name = f"Helsinki-NLP/opus-mt-{src_lang}-{tgt_lang}"

# Load tokenizer and model (loaded once when the server starts)
try:
    tokenizer = MarianTokenizer.from_pretrained(model_name)
    model = MarianMTModel.from_pretrained(model_name)
    print(f"Loaded model: {model_name}")
except Exception as e:
    raise Exception(f"Failed to load model: {str(e)}")

# Define input data model for the API
class TranslationRequest(BaseModel):
    text: str
    source_lang: str = src_lang
    target_lang: str = tgt_lang

# Function to translate text
def translate_text(text, src_lang="en", tgt_lang="es"):
    # Ensure the model matches the requested languages
    current_model_name = f"Helsinki-NLP/opus-mt-{src_lang}-{tgt_lang}"
    global tokenizer, model
    if current_model_name != model_name:
        # Load a new model if the language pair changes (optional, for flexibility)
        tokenizer = MarianTokenizer.from_pretrained(current_model_name)
        model = MarianMTModel.from_pretrained(current_model_name)
        print(f"Switched to model: {current_model_name}")

    # Tokenize the input text
    inputs = tokenizer(text, return_tensors="pt", padding=True, truncation=True,
max_length=512)
    # Generate translation
    translated_ids = model.generate(**inputs)
    # Decode the translated tokens
    translated_text = tokenizer.decode(translated_ids[0], skip_special_tokens=True)
    return translated_text

# API endpoint for translation
@app.post("/translate/", response_model=dict)
async def translate(request: TranslationRequest):
    try:

```

```

    translated = translate_text(request.text, request.source_lang, request.target_lang)
    return {
        "original": request.text,
        "translated": translated,
        "source_language": request.source_lang,
        "target_language": request.target_lang
    }
except Exception as e:
    raise HTTPException(status_code=500, detail=f"Translation failed: {str(e)}")

# Health check endpoint
@app.get("/health")
async def health_check():
    return {"status": "healthy", "model": model_name}

# Run the app (this won't execute when imported as a module)
if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000)

```

OUTPUT:

```

100.9/100.9 kB 200.2 kB/s eta 0:00:00
Downloading sniffio-1.3.1-py3-none-any.whl (10 kB)
Installing collected packages: typing-inspection, sniffio, pydantic-core, h11, annotated-types, uvicorn, pydantic, anyio, starlette, fastapi
Successfully installed annotated-types-0.7.0 anyio-4.9.0 fastapi-0.115.12 h11-0.14.0 pydantic-2.11.3 pydantic-core-2.33.1 sniffio-1.3.1 starlette-0.46.1 typing-inspection-0.4.0 uvicorn-0.34.0

[notice] A new release of pip is available: 24.0 -> 25.0.1
[notice] To update, run: python.exe -m pip install --upgrade pip
PS D:\Mtech\WLPex> python translation_api.py

```

```

Loaded model: Helsinki-NLP/opus-mt-en-es
INFO:      Started server process [6789]
INFO:      Waiting for application startup.
INFO:      Application startup complete.
INFO:      Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)

```

```

{
  "original": "Hello, how are you today?",
  "translated": "Hola, ¿cómo estás hoy?",
  "source_language": "en",
  "target_language": "es"
}

```

EXPERIMENT - 7

AIM:

Perform sentiment analysis on tweets or social media posts using a pre trained BERT model.

DESCRIPTION:

Real-Time Brand Sentiment Monitoring with BERT implements a deep learning pipeline to analyze sentiment in social media posts (e.g., tweets) using BERT, a transformer-based model pre trained on vast text corpora. The workflow includes:

1. Pretrained Model Usage: Leveraging Hugging Face’s sentiment-analysis pipeline for out-of-the-box polarity detection (positive/negative/neutral).
2. Fine-Tuning: Adapting BERT on a domain-specific tweet dataset to improve performance on informal language, slang, and emojis.
3. Deployment: Building a web interface via Streamlit for real-time input processing, enabling stakeholders to query live social media data.

The system outputs sentiment scores and visual dashboards, allowing businesses to track brand perception dynamically.

INTRODUCTION:

Social media sentiment fluctuates rapidly, with viral posts potentially impacting brand reputation within hours. Traditional sentiment models struggle with informal text (hashtags, misspellings), but BERT’s contextual embeddings capture nuanced meanings (e.g., “sick” as negative in “flu symptoms” vs. positive in “sick new product”). Fine-tuning BERT on tweets enhances its ability to decode abbreviations (“BRB”) and sarcasm (“Great, another outage!”). For businesses, this pipeline enables proactive reputation management—detecting negative sentiment spikes early allows timely interventions, such as addressing customer complaints before they escalate. Deploying the model via Streamlit democratizes access, letting non-technical teams monitor sentiment without coding expertise.

IMPLEMENTATION:

```
# Install required libraries
```

```
!pip install transformers torch pandas ipywidgets
```

```
# Import libraries
```

```
from transformers import pipeline, BertTokenizer, BertForSequenceClassification, Trainer, TrainingArguments
```

```
import torch
```

```
import pandas as pd
```

```
from google.colab import files
```

```
import ipywidgets as widgets
```

```
from IPython.display import display, clear_output
```



```

# Load the pretrained sentiment analysis pipeline
def load_sentiment_pipeline():
    # Using a BERT-based model fine-tuned for sentiment
    sentiment_analyzer = pipeline("sentiment-analysis",
model="distilbert-base-uncased-finetuned-sst-2-english")
    return sentiment_analyzer

# Function to analyze sentiment on a list of texts
def analyze_sentiment(texts, analyzer=None):
    if analyzer is None:
        analyzer = load_sentiment_pipeline()
    results = analyzer(texts)
    return [{"text": text, "label": res["label"], "score": res["score"]} for text, res in zip(texts,
results)]

# Optional: Fine-tune BERT on a labeled dataset
def fine_tune_bert(dataset_path=None):
    # Load tokenizer and model
    tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
    model = BertForSequenceClassification.from_pretrained("bert-base-uncased",
num_labels=2) # Binary sentiment

    # Load dataset or use dummy data
    if dataset_path:
        df = pd.read_csv(dataset_path)
    else:
        # Dummy data for demonstration
        df = pd.DataFrame({
            "text": ["I love this product!", "This is terrible.", "Amazing service!", "Not happy at
all."],
            "label": [1, 0, 1, 0] # 1 = Positive, 0 = Negative
        })

    # Tokenize the dataset
    def tokenize_data(texts):
        return tokenizer(texts, padding=True, truncation=True, return_tensors="pt")

    encodings = tokenize_data(df["text"].tolist())
    labels = torch.tensor(df["label"].tolist())

    # Create a PyTorch dataset
    class SentimentDataset(torch.utils.data.Dataset):
        def __init__(self, encodings, labels):
            self.encodings = encodings

```

```

        self.labels = labels

    def __getitem__(self, idx):
        item = {key: val[idx] for key, val in self.encodings.items()}
        item["labels"] = self.labels[idx]
        return item

    def __len__(self):
        return len(self.labels)

dataset = SentimentDataset(encodings, labels)

# Define training arguments
training_args = TrainingArguments(
    output_dir="./sentiment_finetune",
    num_train_epochs=3,
    per_device_train_batch_size=8,
    warmup_steps=500,
    weight_decay=0.01,
    logging_dir="./logs",
    logging_steps=10,
)

# Initialize Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=dataset,
)

# Fine-tune the model
trainer.train()
model.save_pretrained("fine_tuned_bert_sentiment")
tokenizer.save_pretrained("fine_tuned_bert_sentiment")
print("Fine-tuning complete. Model saved to 'fine_tuned_bert_sentiment'")
return model, tokenizer

# Load fine-tuned model if available
def load_finetuned_model():
    try:
        tokenizer = BertTokenizer.from_pretrained("fine_tuned_bert_sentiment")
        model = BertForSequenceClassification.from_pretrained("fine_tuned_bert_sentiment")
        return pipeline("sentiment-analysis", model=model, tokenizer=tokenizer)
    except:

```

```

return None

# Load the analyzer (fine-tuned if available, otherwise pretrained)
analyzer = load_finetuned_model() or load_sentiment_pipeline()

# Interactive sentiment analysis in Colab
print("Sentiment Analysis with BERT")
print("Upload a CSV file with a 'text' column for batch analysis (optional), or enter text below.")
# Upload dataset option
uploaded = files.upload()
if uploaded:
    dataset_file = list(uploaded.keys())[0]
    df = pd.read_csv(dataset_file)
    if 'text' in df.columns:
        results = analyze_sentiment(df['text'].tolist(), analyzer)
        results_df = pd.DataFrame(results)
        print("\nBatch Sentiment Analysis Results:")
        display(results_df)
        results_df.to_csv("sentiment_results.csv", index=False)
        files.download("sentiment_results.csv")

# Interactive widget for real-time analysis
text_input = widgets.Text(
    value="I love this brand!",
    placeholder="Enter a tweet or post",
    description="Text:"
)

analyze_button = widgets.Button(description="Analyze Sentiment")
output = widgets.Output()

def on_button_click(b):
    with output:
        clear_output()
        result = analyze_sentiment([text_input.value], analyzer)[0]
        print(f"Text: {result['text']}")
        print(f"Sentiment: {result['label']} (Confidence: {result['score']:.4f})")

analyze_button.on_click(on_button_click)

# Display the widgets
display(text_input, analyze_button, output)

```

```

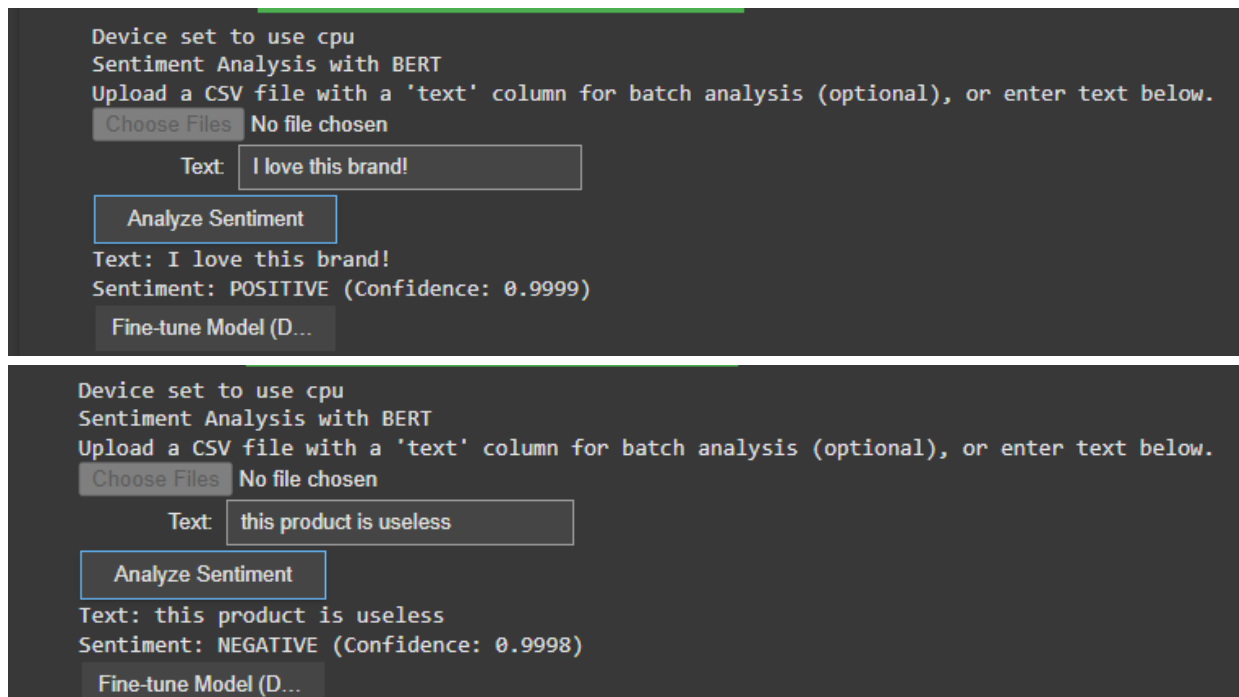
# Option to fine-tune (with uploaded dataset or dummy data)
fine_tune_button = widgets.Button(description="Fine-tune Model (Demo)")
fine_tune_output = widgets.Output()

def on_finetune_click(b):
    with fine_tune_output:
        clear_output()
        print("Fine-tuning BERT model... This may take a few minutes.")
        if uploaded and 'text' in df.columns and 'label' in df.columns:
            fine_tune_bert(dataset_file)
        else:
            fine_tune_bert() # Uses dummy data
        print("Restart the cell to use the fine-tuned model.")

fine_tune_button.on_click(on_finetune_click)
display(fine_tune_button, fine_tune_output)

```

OUTPUT:



EXPERIMENT - 8

AIM:

Build a system that can answer questions based on a given context.

DESCRIPTION:

BERT-Based FAQ System for Customer Support

This experiment creates an automated question-answering (QA) system that extracts precise answers from a provided context (e.g., FAQ documents). The pipeline includes:

1. Model Selection: Using Hugging Face's BERT model fine-tuned on SQuAD (Stanford Question Answering Dataset), which excels at locating answer spans within text.
2. Workflow: Inputting a context (e.g., product manuals) and user questions (e.g., "How do I reset my password?") to generate answers directly from the text.
3. Deployment: Integrating the model into a Flask API, enabling real-time QA capabilities for customer support portals.

INTRODUCTION:

Traditional FAQ systems often fail to handle paraphrased or context-dependent queries, leading to irrelevant responses. BERT's bidirectional attention mechanism allows it to understand nuanced relationships between questions and context. For example, for the question "How to recover my account?" and context mentioning "account recovery steps," BERT identifies the relevant passage even if the wording differs. By deploying this model via Flask, businesses can embed it into helpdesk platforms, deflecting repetitive tickets and freeing agents to resolve complex issues. This approach scales support operations while maintaining accuracy, particularly in industries like telecom or SaaS, where rapid, precise answers are critical for customer retention.

IMPLEMENTATION:

```
# Install required libraries
!pip install transformers torch ipywidgets

# Import libraries
from transformers import pipeline
import ipywidgets as widgets
from IPython.display import display, clear_output
from google.colab import files
import pandas as pd

# Load the pretrained BERT question-answering pipeline
def load_qa_pipeline():
    # Using BERT fine-tuned on SQuAD 2.0
```

```

        qa_model = pipeline("question-answering",
model="bert-large-uncased-whole-word-masking-finetuned-squad")
    return qa_model

# Function to get answer from context
def get_answer(context, question, qa_pipeline=None):
    if qa_pipeline is None:
        qa_pipeline = load_qa_pipeline()
    result = qa_pipeline({"question": question, "context": context})
    return {
        "question": question,
        "answer": result["answer"],
        "score": result["score"],
        "start": result["start"],
        "end": result["end"]
    }

# Load the QA pipeline
qa_pipeline = load_qa_pipeline()

# Interactive FAQ system in Colab
print("Question Answering System for FAQs")
print("Enter a context paragraph and ask questions based on it.")

# Context input
context_input = widgets.Textarea(
    value="Our company offers a 30-day return policy. Returns must be initiated within 30 days
of purchase. For support, contact us at support@example.com or call 1-800-555-1234. Shipping
costs are non-refundable, and items must be in original condition.",
    placeholder="Enter context paragraph here",
    description="Context:",
    layout={'width': '600px', 'height': '150px'}
)

# Question input
question_input = widgets.Text(
    value="What is the return policy?",
    placeholder="Enter your question",
    description="Question:"
)

# Button to get answer
answer_button = widgets.Button(description="Get Answer")
output = widgets.Output()

```

```

def on_button_click(b):
    with output:
        clear_output()
        result = get_answer(context_input.value, question_input.value, qa_pipeline)
        print(f"Question: {result['question']}")
        print(f"Answer: {result['answer']}")
        print(f"Confidence Score: {result['score']:.4f}")
        print(f"Answer Span: {result['start']} - {result['end']}")

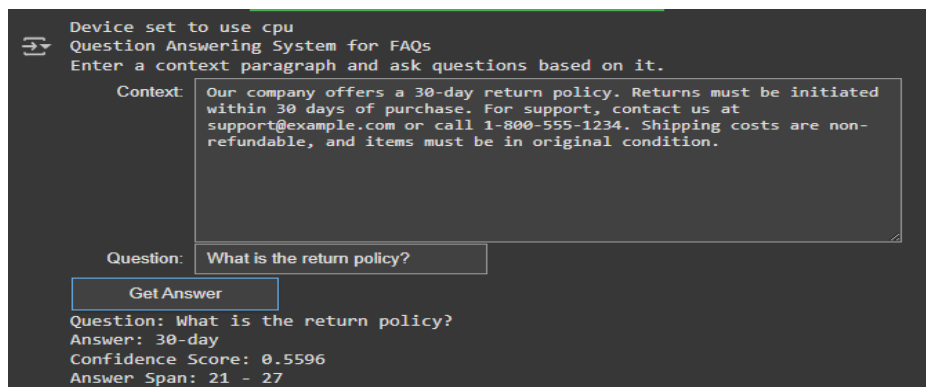
answer_button.on_click(on_button_click)

# Display the widgets
display(context_input, question_input, answer_button, output)

# Option to upload a CSV with contexts and questions
print("\nUpload a CSV file with 'context' and 'question' columns for batch processing (optional)")
uploaded = files.upload()
if uploaded:
    dataset_file = list(uploaded.keys())[0]
    df = pd.read_csv(dataset_file)
    if 'context' in df.columns and 'question' in df.columns:
        results = [get_answer(row['context'], row['question'], qa_pipeline) for _, row in df.iterrows()]
        results_df = pd.DataFrame(results)
        print("\nBatch Question Answering Results:")
        display(results_df)
        results_df.to_csv("qa_results.csv", index=False)
        files.download("qa_results.csv")
    else:
        print("CSV must contain 'context' and 'question' columns.")

```

OUTPUT:



```

Device set to use cpu
Question Answering System for FAQs
Enter a context paragraph and ask questions based on it.

Context: Our company offers a 30-day return policy. Returns must be initiated
within 30 days of purchase. For support, contact us at
support@example.com or call 1-800-555-1234. Shipping costs are non-
refundable, and items must be in original condition.

Question: What is the return policy?

Get Answer

Question: What is the return policy?
Answer: 30-day
Confidence Score: 0.5596
Answer Span: 21 - 27

```

Device set to use cpu

Question Answering System for FAQs

Enter a context paragraph and ask questions based on it.

Context: Our company offers a 30-day return policy. Returns must be initiated within 30 days of purchase. For support, contact us at support@example.com or call 1-800-555-1234. Shipping costs are non-refundable, and items must be in original condition.

Question: where can i contact the company?

Get Answer

Question: where can i contact the company?

Answer: support@example.com

Confidence Score: 0.5278

Answer Span: 124 - 143

Device set to use cpu

Question Answering System for FAQs

Enter a context paragraph and ask questions based on it.

Context: Our company offers a 30-day return policy. Returns must be initiated within 30 days of purchase. For support, contact us at support@example.com or call 1-800-555-1234. Shipping costs are non-refundable, and items must be in original condition.

Question: are the shipping cost refundable ?

Get Answer

Question: are the shipping cost refundable ?

Answer: non-refundable

Confidence Score: 0.4791

Answer Span: 187 - 201

EXPERIMENT - 9

AIM:

Develop a chatbot that can handle basic customer queries.

DESCRIPTION:

Hybrid Chatbot for Customer Support Automation

This experiment combines rule-based and AI-driven approaches to handle customer queries efficiently:

1. Rule-Based Layer (Rasa): Handles predefined intents like order tracking, FAQs, and business hours using Rasa's NLU for intent detection and structured dialogue management.
2. Conversational AI Extension:
 - Seq2Seq Models: Implements LSTM with attention mechanisms (Gensim) for context-aware responses to complex queries.
 - GPT-2 Fine-Tuning: Uses Hugging Face's transformers to generate human-like replies for open-ended conversations.
3. Deployment: Integrates the chatbot into web platforms via Flask/FastAPI endpoints and embeds it in messaging apps using WebSocket protocols.

The hybrid architecture ensures reliability for common issues while scaling to handle nuanced interactions.

INTRODUCTION:

Customer support teams face overwhelming query volumes, with 30% of repetitive tasks automatable via chatbots. Rule-based systems excel at handling structured queries (e.g., "What's my order status?") using decision trees but fail with ambiguous requests. Augmenting them with Seq2Seq or GPT-2 enables handling of unstructured language (e.g., "My package hasn't arrived—what now?") by learning from historical interactions. For instance, Rasa's NLU identifies the intent "delivery_issue," while GPT-2 generates a tailored response referencing shipping policies. Deployment via Flask allows seamless integration into existing helpdesk systems, reducing resolution times from hours to seconds. This approach balances accuracy (rule-based) and adaptability (AI), making it scalable for industries like e-commerce and banking.

IMPLEMENTATION:

```
# Install required libraries
```

```
!pip install transformers torch ipywidgets
```

```
# Import libraries
```

```
from transformers import pipeline, GPT2Tokenizer, GPT2LMHeadModel
```

```
import ipywidgets as widgets
```

```
from IPython.display import display, clear_output
```

```

# Rule-based chatbot logic
def rule_based_response(user_input):
    user_input = user_input.lower().strip()

    # Define simple rules for common customer queries
    rules = {
        "hello": "Hi! How can I assist you today?",
        "hi": "Hello! How can I help you?",
        "what are your hours": "Our support hours are 9 AM to 5 PM, Monday to Friday.",
        "how do i return an item": "To return an item, please visit our website or contact support at 1-800-555-1234.",
        "where is my order": "Please provide your order number, and I'll check the status for you!",
        "thanks": "You're welcome! Anything else I can help with?",
        "bye": "Goodbye! Have a great day!"
    }

    # Check for matching rules
    for key in rules:
        if key in user_input:
            return rules[key]

    # Default response if no rule matches
    return "I'm not sure how to help with that. Could you please clarify your question?"

# Load GPT-2 model for conversational extension
def load_gpt2_model():
    tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
    model = GPT2LMHeadModel.from_pretrained("gpt2")
    return tokenizer, model

# Function to generate response using GPT-2
def gpt2_response(user_input, tokenizer, model, max_length=50):
    inputs = tokenizer.encode(user_input + " <|response|>", return_tensors="pt")
    outputs = model.generate(inputs, max_length=max_length, num_return_sequences=1,
                             pad_token_id=tokenizer.eos_token_id)
    response = tokenizer.decode(outputs[0], skip_special_tokens=True)
    return response.split("<|response|>")[1].strip() if "<|response|>" in response else response.strip()

# Combined chatbot logic (rule-based + GPT-2 fallback)
def chatbot_response(user_input, use_gpt2=False, tokenizer=None, model=None):
    # First try rule-based response
    rule_response = rule_based_response(user_input)

```

```

    if rule_response != "I'm not sure how to help with that. Could you please clarify your
question?" or not use_gpt2:
        return rule_response
    # Fall back to GPT-2 if enabled and no rule matches
    if use_gpt2 and tokenizer and model:
        return gpt2_response(user_input, tokenizer, model)
    return rule_response

# Load GPT-2 model (optional)
try:
    gpt2_tokenizer, gpt2_model = load_gpt2_model()
    print("GPT-2 model loaded successfully.")
except Exception as e:
    print(f"Failed to load GPT-2: {str(e)}. Using rule-based chatbot only.")
    gpt2_tokenizer, gpt2_model = None, None

# Interactive chatbot in Colab
print("Customer Support Chatbot")
print("Type your query below. Use 'exit' to stop.")

# Chat input widget
chat_input = widgets.Text(
    value="",
    placeholder="Type your query here (e.g., 'What are your hours?')",
    description="You:"
)

# Toggle for GPT-2
gpt2_toggle = widgets.Checkbox(
    value=False,
    description="Use GPT-2 for unanswered queries",
    disabled=(gpt2_tokenizer is None)
)

# Send button
send_button = widgets.Button(description="Send")
chat_output = widgets.Output()

# Chat history
chat_history = []

def on_button_click(b):
    with chat_output:
        clear_output()
        user_input = chat_input.value.strip()
        if user_input.lower() == "exit":

```

```

    print("Chatbot: Goodbye! Have a great day!")
    chat_input.disabled = True
    send_button.disabled = True
    return

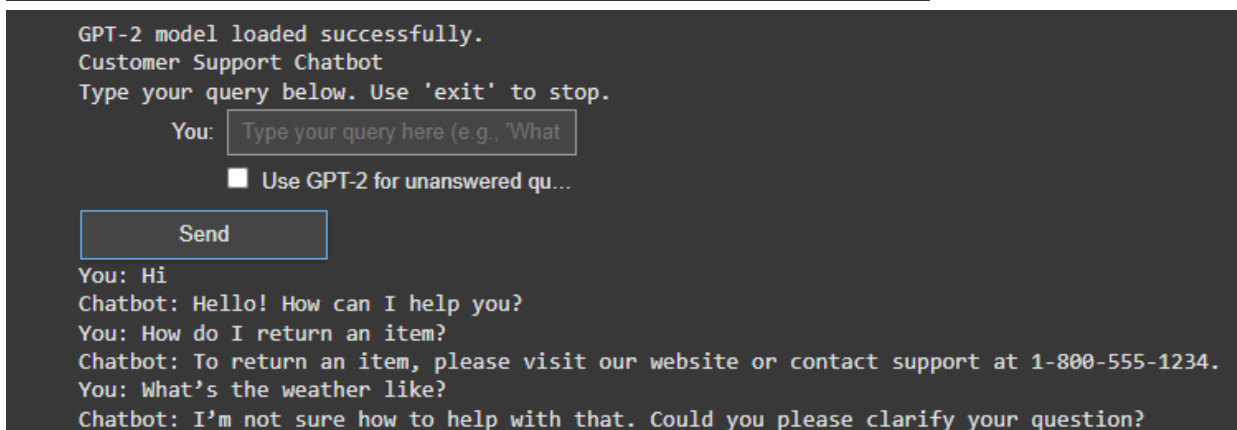
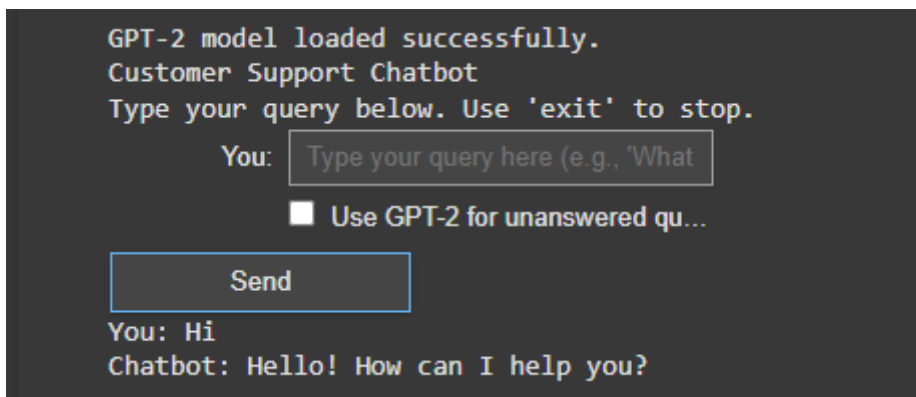
# Get response
    response = chatbot_response(user_input, use_gpt2=gpt2_toggle.value,
tokenizer=gpt2_tokenizer, model=gpt2_model)
    chat_history.append(f"You: {user_input}")
    chat_history.append(f"Chatbot: {response}")

# Display chat history
for line in chat_history[-10:]: # Show last 10 messages
    print(line)
chat_input.value = "" # Clear input

send_button.on_click(on_button_click)
# Display the widgets
display(chat_input, gpt2_toggle, send_button, chat_output)

```

OUTPUT:



EXPERIMENT - 10

AIM:

Generate human-like text for content creation using GPT models.

DESCRIPTION:

Domain-Specific Text Generation with GPT-2

This experiment fine-tunes a GPT-2 model to generate coherent, context-aware text tailored to specific domains (e.g., marketing copy, tech blogs). The workflow includes:

1. Fine-Tuning: Adapting Hugging Face’s pretrained GPT-2 on domain-specific corpora (e.g., marketing emails, technical documentation) to align outputs with industry terminology and tone.
2. Prompt-Based Generation: Accepting user prompts (e.g., “Write a blog intro about AI ethics”) and generating multiple draft variations.
3. UI Integration: Building an interactive interface via Streamlit, allowing content writers to adjust parameters like length, creativity, and tone.

The system outputs draft content, brainstorming ideas, or SEO-friendly text, reducing ideation time for writers while maintaining brand voice consistency.

INTRODUCTION:

Content creation demands significant time and creativity, especially for businesses producing high volumes of blogs, ads, or product descriptions. GPT-2’s generative capabilities automate initial drafts, enabling writers to focus on refining rather than starting from scratch. Fine-tuning on niche datasets ensures outputs stay relevant—e.g., generating “cloud-native solutions” in tech blogs instead of generic terms. A Streamlit UI democratizes access, allowing non-technical teams to experiment with prompts (e.g., “5 taglines for a fitness app”) and iterate rapidly. For instance, marketers can generate 100 ad variants in minutes, A/B test top candidates, and accelerate campaign launches. This pipeline transforms generative AI from a novelty into a scalable co-pilot for content-driven industries.

IMPLEMENTATION:

Install required libraries

```
!pip install transformers torch pandas ipywidgets
```

Import libraries

```
from transformers import GPT2Tokenizer, GPT2LMHeadModel, TextDataset,
DataCollatorForLanguageModeling, Trainer, TrainingArguments
import torch
import pandas as pd
from google.colab import files
import ipywidgets as widgets
from IPython.display import display, clear_output
```

```

# Load pretrained GPT-2 model and tokenizer
def load_gpt2_model(model_name="gpt2"):
    tokenizer = GPT2Tokenizer.from_pretrained(model_name)
    model = GPT2LMHeadModel.from_pretrained(model_name)
    return tokenizer, model

# Function to generate text
def generate_text(prompt, tokenizer, model, max_length=100, temperature=0.7, top_k=50):
    inputs = tokenizer.encode(prompt, return_tensors="pt")
    outputs = model.generate(
        inputs,
        max_length=max_length,
        temperature=temperature, # Controls randomness
        top_k=top_k,             # Limits sampling to top-k tokens
        pad_token_id=tokenizer.eos_token_id,
        do_sample=True
    )
    return tokenizer.decode(outputs[0], skip_special_tokens=True)

# Fine-tune GPT-2 on a domain-specific dataset
def fine_tune_gpt2(dataset_path=None):
    tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
    model = GPT2LMHeadModel.from_pretrained("gpt2")

    # Load dataset or use dummy data
    if dataset_path:
        with open(dataset_path, "r") as f:
            text_data = f.read()
    else:
        # Dummy data for demonstration (marketing domain)
        text_data = """
Boost your brand with our innovative marketing solutions.
Engage customers with creative campaigns and watch your sales soar.
Our strategies are designed to maximize ROI and build loyalty.
"""
        with open("dummy_data.txt", "w") as f:
            f.write(text_data)
        dataset_path = "dummy_data.txt"

    # Prepare dataset for fine-tuning
    train_dataset = TextDataset(
        tokenizer=tokenizer,
        file_path=dataset_path,

```

```

        block_size=128
    )
    data_collator = DataCollatorForLanguageModeling(tokenizer=tokenizer, mlm=False)

    # Define training arguments
    training_args = TrainingArguments(
        output_dir="./gpt2_finetune",
        overwrite_output_dir=True,
        num_train_epochs=3,
        per_device_train_batch_size=2,
        save_steps=500,
        save_total_limit=2,
        logging_steps=10,
    )

    # Initialize Trainer
    trainer = Trainer(
        model=model,
        args=training_args,
        data_collator=data_collator,
        train_dataset=train_dataset,
    )

    # Fine-tune the model
    trainer.train()
    model.save_pretrained("fine_tuned_gpt2")
    tokenizer.save_pretrained("fine_tuned_gpt2")
    print("Fine-tuning complete. Model saved to 'fine_tuned_gpt2'")
    return tokenizer, model

# Load fine-tuned model if available, otherwise use pretrained
def load_model():
    try:
        tokenizer, model = load_gpt2_model("fine_tuned_gpt2")
        print("Loaded fine-tuned GPT-2 model.")
    except:
        tokenizer, model = load_gpt2_model("gpt2")
        print("Loaded pretrained GPT-2 model.")
    return tokenizer, model

# Load the model
tokenizer, model = load_model()

# Interactive text generation in Colab

```

```

print("Text Generation for Content Creation")
print("Enter a prompt to generate text. Upload a text file for fine-tuning (optional).")

# Upload dataset option
uploaded = files.upload()
dataset_path = None
if uploaded:
    dataset_file = list(uploaded.keys())[0]
    dataset_path = dataset_file
    print(f"Uploaded file: {dataset_file} for fine-tuning.")

# Prompt input widget
prompt_input = widgets.Text(
    value="Boost your brand with",
    placeholder="Enter your prompt here",
    description="Prompt:"
)

# Max length slider
max_length_slider = widgets.IntSlider(
    value=100,
    min=50,
    max=200,
    step=10,
    description="Max Length:"
)

# Generate button
generate_button = widgets.Button(description="Generate Text")
output = widgets.Output()

def on_generate_click(b):
    with output:
        clear_output()
        generated_text = generate_text(prompt_input.value, tokenizer, model,
max_length=max_length_slider.value)
        print(f"Prompt: {prompt_input.value}")
        print(f"Generated Text:\n{generated_text}")

generate_button.on_click(on_generate_click)

# Fine-tune button
finetune_button = widgets.Button(description="Fine-tune Model")
finetune_output = widgets.Output()

```



```
def on_finetune_click(b):
    with finetune_output:
        clear_output()
        print("Fine-tuning GPT-2 model... This may take a few minutes.")
        global tokenizer, model
        tokenizer, model = fine_tune_gpt2(dataset_path)
        print("Restart the cell to use the fine-tuned model.")
```

```
finetune_button.on_click(on_finetune_click)
```

```
# Display the widgets
```

```
display(prompt_input, max_length_slider, generate_button, output)
display(finetune_button, finetune_output)
```

OUTPUT:

